

References

Complexity-Accuracy Tradeoffs in RAG: A Comparative Study on HotpotQA

Ghina Al Shdaifat, Huizhen Jin, Shengduo Li, Yixuan Wang, Judy Yang

- [1] Amazon Web Services. *What is Retrieval-Augmented Generation?* <https://aws.amazon.com/what-is/retrieval-augmented-generation/>. Accessed: 2024-12-03. 2024.
- [2] Akari Asai et al. “Self-RAG: Learning to Retrieve, Generate, and Critique Through Self-Reflection”. In: *arXiv preprint arXiv:2310.11511* (2023).
- [3] Patrick Lewis et al. “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”. In: *arXiv preprint arXiv:2005.11401* (2020).
- [4] Chaitanya Sharma. “Retrieval-Augmented Generation: A Comprehensive Survey of Architectures, Enhancements, and Robustness Frontiers”. In: *arXiv preprint arXiv:2506.00054* (2025).
- [5] Hugo Touvron et al. “Llama 2: Open foundation and fine-tuned chat models”. In: *arXiv preprint arXiv:2307.09288* (2023).
- [6] Sentence Transformers. *all-MiniLM-L6-v2*. <https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2>. 2021.
- [7] Zhepei Wei, Wei-Lin Chen, and Yu Meng. “Instructrag: Instructing retrieval-augmented generation via self-synthesized rationales”. In: *arXiv preprint arXiv:2406.13629* (2024).
- [8] Zhilin Yang et al. “HotpotQA: A dataset for diverse, explainable multi-hop question answering”. In: *arXiv preprint arXiv:1809.09600* (2018).